

✓ Homework 4

Instructions

- This homework focuses on understanding and applying CoCoOp for CLIP prompt tuning. It consists of **four questions** designed to assess both theoretical understanding and practical application.
- Please organize your answers and results for the questions below and submit this jupyter notebook as a **.pdf file**.
- **Deadline: 11/26 (Sat) 23:59**

✓ Preparation

- Run the code below before proceeding with the homework.
- If an error occurs, click 'Run Session Again' and then restart the runtime from the beginning.

```
!git clone https://github.com/mlvlab/ProMetaR.git
%cd ProMetaR/

!git clone https://github.com/KaiyangZhou/Dassl.pytorch.git
%cd Dassl.pytorch/

# Install dependencies
!pip install -r requirements.txt
!cp -r dassl ../
# Install this library (no need to re-build if the source code is modified)
# !python setup.py develop
%cd ..

!pip install -r requirements.txt

%mkdir outputs
%mkdir data

%cd data
%mkdir eurosat
!wget http://madm.dfki.de/files/sentinel/EuroSAT.zip -O EuroSAT.zip

!unzip -o EuroSAT.zip -d eurosat/
%cd eurosat
!gdown 1Ip7yaCWF0ea0FUGga0lUdVi_DDQth1o

%cd ../../

import os.path as osp
from collections import OrderedDict
import math
import torch
import torch.nn as nn
from torch.nn import functional as F
from torch.cuda.amp import GradScaler, autocast
from PIL import Image
import torchvision.transforms as transforms
import torch
from clip import clip
from clip.simple_tokenizer import SimpleTokenizer as _Tokenizer
import time
from tqdm import tqdm
import datetime
import argparse
from dassl.utils import setup_logger, set_random_seed, collect_env_info
from dassl.config import get_cfg_default
from dassl.engine import build_trainer
```

```

from dassl.engine import TRAINER_REGISTRY, TrainerX
from dassl.metrics import compute_accuracy
from dassl.utils import load_pretrained_weights, load_checkpoint
from dassl.optim import build_optimizer, build_lr_scheduler

# custom
import datasets.oxford_pets
import datasets.oxford_flowers
import datasets.fgvc_aircraft
import datasets.dtd
import datasets.eurosat
import datasets.stanford_cars
import datasets.food101
import datasets.sun397
import datasets.caltech101
import datasets.ucf101
import datasets.imagenet
import datasets.imagenet_sketch
import datasets.imagenetv2
import datasets.imagenet_a
import datasets.imagenet_r

def print_args(args, cfg):
    print("*****")
    print("** Arguments **")
    print("*****")
    optkeys = list(args.__dict__.keys())
    optkeys.sort()
    for key in optkeys:
        print("{}: {}".format(key, args.__dict__[key]))
    print("*****")
    print("** Config **")
    print("*****")
    print(cfg)

def reset_cfg(cfg, args):
    if args.root:
        cfg.DATASET.ROOT = args.root
    if args.output_dir:
        cfg.OUTPUT_DIR = args.output_dir
    if args.seed:
        cfg.SEED = args.seed
    if args.trainer:
        cfg.TRAINER.NAME = args.trainer
    cfg.DATASET.NUM_SHOTS = 16
    cfg.DATASET.SUBSAMPLE_CLASSES = args.subsample_classes
    cfg.DATALOADER.TRAIN_X.BATCH_SIZE = args.train_batch_size
    cfg.OPTIM.MAX_EPOCH = args.epoch

def extend_cfg(cfg):
    """
    Add new config variables.
    """
    from yacs.config import CfgNode as CN
    cfg.TRAINER.COOP = CN()
    cfg.TRAINER.COOP.N_CTX = 16 # number of context vectors
    cfg.TRAINER.COOP.CSC = False # class-specific context
    cfg.TRAINER.COOP.CTX_INIT = "" # initialization words
    cfg.TRAINER.COOP.PREC = "fp16" # fp16, fp32, amp
    cfg.TRAINER.COOP.CLASS_TOKEN_POSITION = "end" # 'middle' or 'end' or 'front'
    cfg.TRAINER.COOP = CN()
    cfg.TRAINER.COOP.N_CTX = 4 # number of context vectors
    cfg.TRAINER.COOP.CTX_INIT = "a photo of a" # initialization words
    cfg.TRAINER.COOP.PREC = "fp16" # fp16, fp32, amp
    cfg.TRAINER.PROMETAR = CN()
    cfg.TRAINER.PROMETAR.N_CTX_VISION = 4 # number of context vectors at the vision branch
    cfg.TRAINER.PROMETAR.N_CTX_TEXT = 4 # number of context vectors at the language branch
    cfg.TRAINER.PROMETAR.CTX_INIT = "a photo of a" # initialization words
    cfg.TRAINER.PROMETAR.PREC = "fp16" # fp16, fp32, amp

```

```

cfg.TRAINER.PROMETAR.PROMPT_DEPTH_VISION = 9 # Max 12, minimum 0, for 0 it will be using shallow :
cfg.TRAINER.PROMETAR.PROMPT_DEPTH_TEXT = 9 # Max 12, minimum 0, for 0 it will be using shallow IVI
cfg.DATASET.SUBSAMPLE_CLASSES = "all" # all, base or new
cfg.TRAINER.PROMETAR.ADAPT_LR = 0.0005
cfg.TRAINER.PROMETAR.LR_RATIO = 0.0005
cfg.TRAINER.PROMETAR.FAST_ADAPTATION = False
cfg.TRAINER.PROMETAR.MIXUP_ALPHA = 0.5
cfg.TRAINER.PROMETAR.MIXUP_BETA = 0.5
cfg.TRAINER.PROMETAR.DIM_RATE=8
cfg.OPTIM_VNET = CN()
cfg.OPTIM_VNET.NAME = "adam"
cfg.OPTIM_VNET.LR = 0.0003
cfg.OPTIM_VNET.WEIGHT_DECAY = 5e-4
cfg.OPTIM_VNET.MOMENTUM = 0.9
cfg.OPTIM_VNET.SGD_DAMPNING = 0
cfg.OPTIM_VNET.SGD_NESTEROV = False
cfg.OPTIM_VNET.RMSPROP_ALPHA = 0.99
cfg.OPTIM_VNET.ADM_BETA1 = 0.9
cfg.OPTIM_VNET.ADM_BETA2 = 0.999
cfg.OPTIM_VNET.STAGED_LR = False
cfg.OPTIM_VNET.NEW_LAYERS = ()
cfg.OPTIM_VNET.BASE_LR_MULT = 0.1
# Learning rate scheduler
cfg.OPTIM_VNET.LR_SCHEDULER = "single_step"
# -1 or 0 means the stepsize is equal to max_epoch
cfg.OPTIM_VNET.STEPSIZE = (-1, )
cfg.OPTIM_VNET.GAMMA = 0.1
cfg.OPTIM_VNET.MAX_EPOCH = 10
# Set WARMUP_EPOCH larger than 0 to activate warmup training
cfg.OPTIM_VNET.WARMUP_EPOCH = -1
# Either linear or constant
cfg.OPTIM_VNET.WARMUP_TYPE = "linear"
# Constant learning rate when type=constant
cfg.OPTIM_VNET.WARMUP_CONS_LR = 1e-5
# Minimum learning rate when type=linear
cfg.OPTIM_VNET.WARMUP_MIN_LR = 1e-5
# Recount epoch for the next scheduler (last_epoch=-1)
# Otherwise last_epoch=warmup_epoch
cfg.OPTIM_VNET.WARMUP_RECOUNT = True

def setup_cfg(args):
    cfg = get_cfg_default()
    extend_cfg(cfg)
    # 1. From the dataset config file
    if args.dataset_config_file:
        cfg.merge_from_file(args.dataset_config_file)
    # 2. From the method config file
    if args.config_file:
        cfg.merge_from_file(args.config_file)
    # 3. From input arguments
    reset_cfg(cfg, args)
    cfg.freeze()
    return cfg

_tokenizer = _Tokenizer()

def load_clip_to_cpu(cfg): # Load CLIP
    backbone_name = cfg.MODEL.BACKBONE.NAME
    url = clip._MODELS[backbone_name]
    model_path = clip._download(url)

    try:
        # loading JIT archive
        model = torch.jit.load(model_path, map_location="cpu").eval()
        state_dict = None

    except RuntimeError:
        state_dict = torch.load(model_path, map_location="cpu")

```

```

if cfg.TRAINER.NAME == "":
    design_trainer = "CoOp"
else:
    design_trainer = cfg.TRAINER.NAME
design_details = {"trainer": design_trainer,
                  "vision_depth": 0,
                  "language_depth": 0, "vision_ctx": 0,
                  "language_ctx": 0}
model = clip.build_model(state_dict or model.state_dict(), design_details)

return model

from dassl.config import get_cfg_default
cfg = get_cfg_default()
cfg.MODEL.BACKBONE.NAME = "ViT-B/16" # Set the vision encoder backbone of CLIP to ViT.
clip_model = load_clip_to_cpu(cfg)

class TextEncoder(nn.Module):
    def __init__(self, clip_model): # 초기화 하는 함수
        super().__init__()
        self.transformer = clip_model.transformer
        self.positional_embedding = clip_model.positional_embedding
        self.ln_final = clip_model.ln_final
        self.text_projection = clip_model.text_projection
        self.dtype = clip_model.dtype

    def forward(self, prompts, tokenized_prompts): # 모델 호출
        x = prompts + self.positional_embedding.type(self.dtype)
        x = x.permute(1, 0, 2) # NLD -> LND
        x = self.transformer(x)
        x = x.permute(1, 0, 2) # LND -> NLD
        x = self.ln_final(x).type(self.dtype)

        # x.shape = [batch_size, n_ctx, transformer.width]
        # take features from the eot embedding (eot_token is the highest number in each sequence)
        x = x[torch.arange(x.shape[0]), tokenized_prompts.argmax(dim=-1)] @ self.text_projection

        return x

@TRAINER_REGISTRY.register(force=True)
class CoCoOp(TrainerX):
    def check_cfg(self, cfg):
        assert cfg.TRAINER.COCOOP.PREC in ["fp16", "fp32", "amp"]

    def build_model(self):
        cfg = self.cfg
        classnames = self.dm.dataset.classnames
        print(f"Loading CLIP (backbone: {cfg.MODEL.BACKBONE.NAME})")
        clip_model = load_clip_to_cpu(cfg)

        if cfg.TRAINER.COCOOP.PREC == "fp32" or cfg.TRAINER.COCOOP.PREC == "amp":
            # CLIP's default precision is fp16
            clip_model.float()

        print("Building custom CLIP")
        self.model = CoCoOpCustomCLIP(cfg, classnames, clip_model)

        print("Turning off gradients in both the image and the text encoder")
        name_to_update = "prompt_learner"

        for name, param in self.model.named_parameters():
            if name_to_update not in name:
                param.requires_grad_(False)

        # Double check
        enabled = set()

```

```

for name, param in self.model.named_parameters():
    if param.requires_grad:
        enabled.add(name)
print(f"Parameters to be updated: {enabled}")

if cfg.MODEL.INIT_WEIGHTS:
    load_pretrained_weights(self.model.prompt_learner, cfg.MODEL.INIT_WEIGHTS)

self.model.to(self.device)
# NOTE: only give prompt_learner to the optimizer
self.optim = build_optimizer(self.model.prompt_learner, cfg.OPTIM)
self.sched = build_lr_scheduler(self.optim, cfg.OPTIM)
self.register_model("prompt_learner", self.model.prompt_learner, self.optim, self.sched)

self.scaler = GradScaler() if cfg.TRAINER.COCOOP.PREC == "amp" else None

# Note that multi-gpu training could be slow because CLIP's size is
# big, which slows down the copy operation in DataParallel
device_count = torch.cuda.device_count()
if device_count > 1:
    print(f"Multiple GPUs detected (n_gpus={device_count}), use all of them!")
    self.model = nn.DataParallel(self.model)

def before_train(self):
    directory = self.cfg.OUTPUT_DIR
    if self.cfg.RESUME:
        directory = self.cfg.RESUME
    self.start_epoch = self.resume_model_if_exist(directory)

    # Remember the starting time (for computing the elapsed time)
    self.time_start = time.time()

def forward_backward(self, batch):
    image, label = self.parse_batch_train(batch)

    model = self.model
    optim = self.optim
    scaler = self.scaler

    prec = self.cfg.TRAINER.COCOOP.PREC
    loss = model(image, label) # Input image 모델 통과
    optim.zero_grad()
    loss.backward() # Backward (역전파)
    optim.step() # 모델 parameter update

    loss_summary = {"loss": loss.item()}

    if (self.batch_idx + 1) == self.num_batches:
        self.update_lr()

    return loss_summary

def parse_batch_train(self, batch):
    input = batch["img"]
    label = batch["label"]
    input = input.to(self.device)
    label = label.to(self.device)
    return input, label

def load_model(self, directory, epoch=None):
    if not directory:
        print("Note that load_model() is skipped as no pretrained model is given")
        return

    names = self.get_model_names()

    # By default, the best model is loaded
    model_file = "model-best.pth.tar"

```

```

    if epoch is not None:
        model_file = "model.pth.tar-" + str(epoch)

    for name in names:
        model_path = osp.join(directory, name, model_file)

        if not osp.exists(model_path):
            raise FileNotFoundError('Model not found at "{}".format(model_path))

        checkpoint = load_checkpoint(model_path)
        state_dict = checkpoint["state_dict"]
        epoch = checkpoint["epoch"]

        # Ignore fixed token vectors
        if "token_prefix" in state_dict:
            del state_dict["token_prefix"]

        if "token_suffix" in state_dict:
            del state_dict["token_suffix"]

        print("Loading weights to {} " 'from "{}' (epoch = {})'.format(name, model_path, epoch))
        # set strict=False
        self._models[name].load_state_dict(state_dict, strict=False)

def after_train(self):
    print("Finish training")

    do_test = not self.cfg.TEST.NO_TEST
    if do_test:
        if self.cfg.TEST.FINAL_MODEL == "best_val":
            print("Deploy the model with the best val performance")
            self.load_model(self.output_dir)
        else:
            print("Deploy the last-epoch model")
        acc = self.test()

    # Show elapsed time
    elapsed = round(time.time() - self.time_start)
    elapsed = str(datetime.timedelta(seconds=elapsed))
    print(f"Elapsed: {elapsed}")

    # Close writer
    self.close_writer()
    return acc

def train(self):
    """Generic training loops."""
    self.before_train()
    for self.epoch in range(self.start_epoch, self.max_epoch):
        self.before_epoch()
        self.run_epoch()
        self.after_epoch()
    acc = self.after_train()
    return acc

parser = argparse.ArgumentParser()
parser.add_argument("--root", type=str, default="data/", help="path to dataset")
parser.add_argument("--output-dir", type=str, default="outputs/cocoo3", help="output directory")
parser.add_argument(
    "--seed", type=int, default=1, help="only positive value enables a fixed seed"
)
parser.add_argument(
    "--config-file", type=str, default="configs/trainers/ProMetaR/vit_b16_c2_ep10_batch4_4+4ctx.yaml",
)
parser.add_argument(
    "--dataset-config-file",
    type=str,
    default="configs/datasets/eurosat.yaml",

```

```

        help="path to config file for dataset setup",
    )
    parser.add_argument("--trainer", type=str, default="CoOp", help="name of trainer")
    parser.add_argument("--eval-only", action="store_true", help="evaluation only")
    parser.add_argument(
        "--model-dir",
        type=str,
        default="",
        help="load model from this directory for eval-only mode",
    )
    parser.add_argument("--train-batch-size", type=int, default=4)
    parser.add_argument("--epoch", type=int, default=10)
    parser.add_argument("--subsample-classes", type=str, default="base")
    parser.add_argument(
        "--load-epoch", type=int, default=0, help="load model weights at this epoch for evaluation"
    )
    args = parser.parse_args([])

def main(args):
    cfg = setup_cfg(args)
    if cfg.SEED >= 0:
        set_random_seed(cfg.SEED)

    if torch.cuda.is_available() and cfg.USE_CUDA:
        torch.backends.cudnn.benchmark = True

    trainer = build_trainer(cfg)
    if args.eval_only:
        trainer.load_model(args.model_dir, epoch=args.load_epoch)
        acc = trainer.test()
        return acc

    acc = trainer.train()
    return acc

```



```
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2150.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1782.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_579.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1025.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2409.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_853.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_421.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_386.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2068.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_882.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_357.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_65.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_736.jpg
/content/ProMetaR/data/eurosat
Downloading...
From: https://drive.google.com/uc?id=1Ip7yaCWFi0ea0FUGga0lUdVi\_DDQ0th1o
To: /content/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
100% 3.01M/3.01M [00:00<00:00, 23.8MB/s]
/content/ProMetaR
100%|██████████| 351M/351M [00:06<00:00, 53.2MiB/s]
```

- Q1. Understanding and implementing CoCoOp

- We have learned how to define CoOp in Lab Session 4.
- The main difference between CoOp and CoCoOp is **meta network** to extract image tokens that is added to the text prompt.
- Based on the CoOp code given in Lab Session 4, fill-in-the-blank exercise (4 blanks!!) to test your understanding of critical parts of the CoCoOp.

```
import torch.nn as nn
```

```

class CoCoOpPromptLearner(nn.Module):
    def __init__(self, cfg, classnames, clip_model):
        super().__init__()
        n_cls = len(classnames)
        n_ctx = cfg.TRAINER.COOCOOP.N_CTX
        ctx_init = cfg.TRAINER.COOCOOP.CTX_INIT
        dtype = clip_model.dtype
        ctx_dim = clip_model.ln_final.weight.shape[0]
        vis_dim = clip_model.visual.output_dim
        clip_imsize = clip_model.visual.input_resolution
        cfg_imsize = cfg.INPUT.SIZE[0]
        assert cfg_imsize == clip_imsize, f"cfg_imsize ({cfg_imsize}) must equal to clip_imsize ({clip_

    if ctx_init:
        # use given words to initialize context vectors
        ctx_init = ctx_init.replace("_", " ")
        n_ctx = len(ctx_init.split(" "))
        prompt = clip.tokenize(ctx_init)
        with torch.no_grad():
            embedding = clip_model.token_embedding(prompt).type(dtype)
            ctx_vectors = embedding[0, 1: 1 + n_ctx, :]
            prompt_prefix = ctx_init
    else:
        # random initialization
        ctx_vectors = torch.empty(n_ctx, ctx_dim, dtype=dtype)
        nn.init.normal_(ctx_vectors, std=0.02)
        prompt_prefix = " ".join(["X"] * n_ctx)

    print(f'Initial context: "{prompt_prefix}"')
    print(f'Number of context words (tokens): {n_ctx}')

    self.ctx = nn.Parameter(ctx_vectors) # Wrap the initialized prompts above as parameters to mal

### Tokenize ###

```



```

classnames = [name.replace("_", " ") for name in classnames] # 예) "Forest"
name_lens = [len(_tokenizer.encode(name)) for name in classnames]
prompts = [prompt_prefix + " " + name + "." for name in classnames] # 예) "A photo of Forest."

tokenized_prompts = torch.cat([clip.tokenize(p) for p in prompts]) # 예) [49406, 320, 1125, 539

```

```

#####
##### Q1. Fill in the blank #####
##### Define Meta Net #####
self.meta_net = nn.Sequential(OrderedDict([
    ("linear1", nn.Linear(vis_dim, vis_dim // 16)),
    ("relu", nn.ReLU(inplace=True)),
    ("linear2", nn.Linear(vis_dim // 16, ctx_dim))
]))
#####
## Hint: meta network is composed to linear layer, relu activation, and linear layer.

```

```

if cfg.TRAINER.COCOOP.PREC == "fp16":
    self.meta_net.half()

with torch.no_grad():
    embedding = clip_model.token_embedding(tokenized_prompts).type(dtype)

# These token vectors will be saved when in save_model(),
# but they should be ignored in load_model() as we want to use
# those computed using the current class names
self.register_buffer("token_prefix", embedding[:, :1, :]) # SOS
self.register_buffer("token_suffix", embedding[:, 1 + n_ctx:, :]) # CLS, EOS
self.n_cls = n_cls
self.n_ctx = n_ctx
self.tokenized_prompts = tokenized_prompts # torch.Tensor
self.name_lens = name_lens

```

```

def construct_prompts(self, ctx, prefix, suffix, label=None):
    # dim0 is either batch_size (during training) or n_cls (during testing)
    # ctx: context tokens, with shape of (dim0, n_ctx, ctx_dim)
    # prefix: the sos token, with shape of (n_cls, 1, ctx_dim)
    # suffix: remaining tokens, with shape of (n_cls, *, ctx_dim)

```

```

    if label is not None:
        prefix = prefix[label]
        suffix = suffix[label]

```

```

    prompts = torch.cat(
        [
            prefix, # (dim0, 1, dim)
            ctx, # (dim0, n_ctx, dim)
            suffix, # (dim0, *, dim)
        ],
        dim=1,
    )

```

```

    return prompts

```

```

def forward(self, im_features):
    prefix = self.token_prefix
    suffix = self.token_suffix
    ctx = self.ctx # (n_ctx, ctx_dim)

```

```

#####
##### Q2,3. Fill in the blank #####
bias = self.meta_net(im_features) # (batch, ctx_dim)
bias = bias.unsqueeze(1) # (batch, 1, ctx_dim)

```

```

ctx = ctx.unsqueeze(0) # (1, n_ctx, ctx_dim)
ctx_shifted = ctx + bias # (batch, n_ctx, ctx_dim)
#####
#####

# Use instance-conditioned context tokens for all classes
prompts = []
for ctx_shifted_i in ctx_shifted:
    ctx_i = ctx_shifted_i.unsqueeze(0).expand(self.n_cls, -1, -1)
    pts_i = self.construct_prompts(ctx_i, prefix, suffix) # (n_cls, n_tkn, ctx_dim)
    prompts.append(pts_i)
prompts = torch.stack(prompts)

return prompts

class CoCoOpCustomCLIP(nn.Module):
    def __init__(self, cfg, classnames, clip_model):
        super().__init__()
        self.prompt_learner = CoCoOpPromptLearner(cfg, classnames, clip_model)
        self.tokenized_prompts = self.prompt_learner.tokenized_prompts
        self.image_encoder = clip_model.visual
        self.text_encoder = TextEncoder(clip_model)
        self.logit_scale = clip_model.logit_scale
        self.dtype = clip_model.dtype

    def forward(self, image, label=None):
        tokenized_prompts = self.tokenized_prompts
        logit_scale = self.logit_scale.exp()

        image_features = self.image_encoder(image.type(self.dtype))
        image_features = image_features / image_features.norm(dim=-1, keepdim=True)

        #####
        ##### Q4. Fill in the blank #####
        prompts = self.prompt_learner(image_features)
        #####
        #####

        logits = []
        for pts_i, imf_i in zip(prompts, image_features):
            text_features = self.text_encoder(pts_i, tokenized_prompts)
            text_features = text_features / text_features.norm(dim=-1, keepdim=True)
            l_i = logit_scale * imf_i @ text_features.t()
            logits.append(l_i)
        logits = torch.stack(logits)

        if self.prompt_learner.training:
            return F.cross_entropy(logits, label)

        return logits

```

✓ Q2. Training CoCoOp

In this task, you will train CoCoOp on the EuroSAT dataset. If your implementation of CoCoOp in Question 1 is correct, the following code should execute without errors. Please submit the execution file so we can evaluate whether your code runs without any issues.

```

# Train on the Base Classes Train split and evaluate accuracy on the Base Classes Test split.
args.trainer = "CoCoOp"
args.train_batch_size = 4
args.epoch = 100

```

```
args.output_dir = "outputs/cocoop"
```

```
args.subsample_classes = "base"
```

```
args.eval_only = False
```

```
cocoop_base_acc = main(args)
```

```
→ epoch [55/100] batch [20/20] time 0.097 (0.124) data 0.000 (0.020) loss 0.2917 (0.1799) lr 1.0933e-
epoch [56/100] batch [20/20] time 0.118 (0.136) data 0.000 (0.020) loss 0.2384 (0.2613) lr 1.0545e-
epoch [57/100] batch [20/20] time 0.135 (0.196) data 0.000 (0.036) loss 0.3364 (0.3352) lr 1.0158e-
epoch [58/100] batch [20/20] time 0.096 (0.127) data 0.000 (0.019) loss 0.3237 (0.2660) lr 9.7732e-
epoch [59/100] batch [20/20] time 0.094 (0.123) data 0.000 (0.015) loss 0.0295 (0.2851) lr 9.3914e-
epoch [60/100] batch [20/20] time 0.094 (0.124) data 0.000 (0.016) loss 0.0961 (0.1896) lr 9.0126e-
epoch [61/100] batch [20/20] time 0.132 (0.138) data 0.000 (0.024) loss 0.3149 (0.2265) lr 8.6373e-
epoch [62/100] batch [20/20] time 0.139 (0.193) data 0.000 (0.029) loss 0.0041 (0.2124) lr 8.2658e-
epoch [63/100] batch [20/20] time 0.104 (0.125) data 0.000 (0.017) loss 0.1748 (0.2624) lr 7.8984e-
epoch [64/100] batch [20/20] time 0.095 (0.132) data 0.000 (0.023) loss 0.2600 (0.1714) lr 7.5357e-
epoch [65/100] batch [20/20] time 0.094 (0.127) data 0.000 (0.018) loss 0.5747 (0.2100) lr 7.1778e-
epoch [66/100] batch [20/20] time 0.120 (0.140) data 0.000 (0.017) loss 0.1279 (0.1686) lr 6.8251e-
epoch [67/100] batch [20/20] time 0.139 (0.192) data 0.000 (0.024) loss 0.0054 (0.2219) lr 6.4781e-
epoch [68/100] batch [20/20] time 0.094 (0.124) data 0.000 (0.015) loss 0.2773 (0.2684) lr 6.1370e-
epoch [69/100] batch [20/20] time 0.094 (0.124) data 0.000 (0.017) loss 0.0228 (0.2471) lr 5.8022e-
epoch [70/100] batch [20/20] time 0.095 (0.123) data 0.000 (0.016) loss 0.2318 (0.1503) lr 5.4740e-
epoch [71/100] batch [20/20] time 0.126 (0.134) data 0.000 (0.017) loss 0.0285 (0.1188) lr 5.1527e-
epoch [72/100] batch [20/20] time 0.139 (0.188) data 0.000 (0.026) loss 0.1163 (0.2144) lr 4.8387e-
epoch [73/100] batch [20/20] time 0.096 (0.126) data 0.000 (0.017) loss 0.0424 (0.1745) lr 4.5322e-
epoch [74/100] batch [20/20] time 0.095 (0.135) data 0.000 (0.017) loss 0.1774 (0.1305) lr 4.2336e-
epoch [75/100] batch [20/20] time 0.093 (0.124) data 0.000 (0.022) loss 0.0523 (0.1880) lr 3.9432e-
epoch [76/100] batch [20/20] time 0.120 (0.137) data 0.000 (0.017) loss 0.0109 (0.1781) lr 3.6612e-
epoch [77/100] batch [20/20] time 0.151 (0.200) data 0.000 (0.029) loss 0.0092 (0.1832) lr 3.3879e-
epoch [78/100] batch [20/20] time 0.093 (0.126) data 0.000 (0.017) loss 0.1420 (0.2149) lr 3.1236e-
epoch [79/100] batch [20/20] time 0.095 (0.125) data 0.000 (0.019) loss 0.6455 (0.2502) lr 2.8686e-
epoch [80/100] batch [20/20] time 0.094 (0.123) data 0.000 (0.022) loss 0.1262 (0.1671) lr 2.6231e-
epoch [81/100] batch [20/20] time 0.124 (0.136) data 0.000 (0.022) loss 0.1049 (0.1736) lr 2.3873e-
epoch [82/100] batch [20/20] time 0.139 (0.199) data 0.000 (0.029) loss 0.5278 (0.1947) lr 2.1615e-
epoch [83/100] batch [20/20] time 0.095 (0.125) data 0.000 (0.016) loss 0.1053 (0.1895) lr 1.9459e-
epoch [84/100] batch [20/20] time 0.105 (0.125) data 0.000 (0.016) loss 0.1261 (0.1526) lr 1.7407e-
epoch [85/100] batch [20/20] time 0.097 (0.124) data 0.000 (0.016) loss 0.0314 (0.1640) lr 1.5462e-
epoch [86/100] batch [20/20] time 0.134 (0.137) data 0.000 (0.021) loss 0.0459 (0.1491) lr 1.3624e-
epoch [87/100] batch [20/20] time 0.137 (0.190) data 0.000 (0.034) loss 0.2108 (0.1862) lr 1.1897e-
epoch [88/100] batch [20/20] time 0.101 (0.124) data 0.000 (0.019) loss 0.1178 (0.2581) lr 1.0281e-
epoch [89/100] batch [20/20] time 0.095 (0.125) data 0.000 (0.023) loss 0.0460 (0.2158) lr 8.7779e-
epoch [90/100] batch [20/20] time 0.091 (0.125) data 0.000 (0.023) loss 0.0492 (0.1039) lr 7.3899e-
epoch [91/100] batch [20/20] time 0.131 (0.137) data 0.000 (0.020) loss 0.2791 (0.1459) lr 6.1179e-
epoch [92/100] batch [20/20] time 0.151 (0.189) data 0.000 (0.033) loss 0.0514 (0.1019) lr 4.9633e-
epoch [93/100] batch [20/20] time 0.100 (0.134) data 0.000 (0.019) loss 0.1763 (0.2449) lr 3.9271e-
epoch [94/100] batch [20/20] time 0.093 (0.124) data 0.000 (0.021) loss 0.2859 (0.2261) lr 3.0104e-
epoch [95/100] batch [20/20] time 0.093 (0.124) data 0.000 (0.015) loss 0.1564 (0.1853) lr 2.2141e-
epoch [96/100] batch [20/20] time 0.135 (0.138) data 0.000 (0.016) loss 0.4089 (0.1330) lr 1.5390e-
epoch [97/100] batch [20/20] time 0.141 (0.187) data 0.000 (0.033) loss 0.0698 (0.1542) lr 9.8566e-
epoch [98/100] batch [20/20] time 0.093 (0.125) data 0.000 (0.018) loss 0.2188 (0.2041) lr 5.5475e-
epoch [99/100] batch [20/20] time 0.095 (0.125) data 0.000 (0.023) loss 0.0691 (0.1264) lr 2.4666e-
epoch [100/100] batch [20/20] time 0.092 (0.125) data 0.000 (0.019) loss 0.0025 (0.1101) lr 6.1680e-
Checkpoint saved to outputs/cocoop/prompt_learner/model.pth.tar-100
```

```
Finish training
```

```
Deploy the last-epoch model
```

```
Evaluate on the *test* set
```

```
100%|██████████| 42/42 [01:05<00:00, 1.57s/it]=> result
```

```
* total: 4,200
```

```
* correct: 3,813
```

```
* accuracy: 90.8%
```

```
* error: 9.2%
```

```
* macro_f1: 90.9%
```

```
Elapsed: 0:06:15
```

```
# Accuracy on the New Classes.
```

```
args.model_dir = "outputs/cocoop"
```

```
args.output_dir = "outputs/cocoop/new_classes"
```

```
args.subsample_classes = "new"
```

```
args.load_epoch = 100
```

```
args.eval_only = True
```

```
coop_novel_acc = main(args)
```

```

➡ Loading trainer: CoCoOp
Loading dataset: EuroSAT
Reading split from /content/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
Loading preprocessed few-shot data from /content/ProMetaR/data/eurosat/split_fewshot/shot_16-seed_1
SUBSAMPLE NEW CLASSES!
Building transform_train
+ random resized crop (size=(224, 224), scale=(0.08, 1.0))
+ random flip
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
Building transform_test
+ resize the smaller edge to 224
+ 224x224 center crop
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
-----
Dataset      EuroSAT
# classes    5
# train_x    80
# val        20
# test       3,900
-----
Loading CLIP (backbone: ViT-B/16)
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose pa
  warnings.warn(
/content/ProMetaR/dassl/utils/torchtools.py:102: FutureWarning: You are using `torch.load` with `we
  checkpoint = torch.load(fpath, map_location=map_location)
Building custom CLIP
Initial context: "a photo of a"
Number of context words (tokens): 4
Turning off gradients in both the image and the text encoder
Parameters to be updated: {'prompt_learner.meta_net.linear2.bias', 'prompt_learner.meta_net.linear1
Loading evaluator: Classification
Loading weights to prompt_learner from "outputs/cocoop/prompt_learner/model.pth.tar-100" (epoch = 1
Evaluate on the *test* set
100%|██████████| 39/39 [00:59<00:00, 1.53s/it]=> result
* total: 3,900
* correct: 1,687
* accuracy: 43.3%
* error: 56.7%
* macro_f1: 39.0%


```

▼ Q3. Analyzing the results of CoCoOp

Compare the results of CoCoOp with those of CoOp that we trained in Lab Session 4. Discuss possible reasons for the performance differences observed between CoCoOp and CoOp.

The result of executing '4_VLM_Prompt_Tuning' shows that CoOp makes 91.40 % accuracy for Base classes test split and 51.46% accuracy for Novel classes which means it has never been trained. The expected outcome was that CoCoOp would perform better. However, the results are oppsite. CoCoOp is designed to allow the input image to guide the learnable prompt by employing a meta-network. This setup regularizes the prompt to prevent it from overfitting to specific classes.

In contrast, in the conventional CoOp, the prompt optimization process relies solely on the text encoder, potentially weakening the transferability of the image representation. When the context tokens are fully optimized for a specific dataset, they can lead to overfitting.



```
import matplotlib.pyplot as plt
import numpy as np

metrics = ['Base', 'Novel']


cocoop_acc_list = [cocoop_base_acc, coop_novel_acc]

bar_width = 0.35
index = np.arange(len(metrics))
fig, ax = plt.subplots()
bar1 = ax.bar(index, cocoop_acc_list, bar_width, label='CoCoOp')

ax.set_ylabel('Scores')
ax.set_title('Model Performance Comparison')
ax.set_xticks(index + bar_width / 2)
ax.set_xticklabels(metrics)
ax.legend()

def add_value_labels(bars):
    for bar in bars:
        height = bar.get_height()
        ax.annotate(f'{height:.2f}', xy=(bar.get_x() + bar.get_width() / 2, height),
                    xytext=(0, 2), # 2 points vertical offset
                    textcoords='offset points',
                    ha='center', va='bottom')

add_value_labels(bar1)
plt.tight_layout()
plt.show()
```



Considering these characteristics, the reasons for CoCoOp's lower performance on novel classes can be explained as follows:

If the EuroSAT novel classes are significantly different from the training classes in terms of distribution, the conditional prompts of CoCoOp might lead to less accurate predictions.

Since CoCoOp is a more complex model, if the dataset's distribution is imbalanced or the dataset is small, the meta-network may not be sufficiently well-trained, resulting in poorer performance on novel datasets.

Regarding the characteristics of the EuroSAT dataset, it is a relatively small dataset. CoCoOp learns class-specific conditional prompts, and in datasets where the class distinctions are clear, the risk of overfitting increases. As a result, while both models may exhibit high accuracy on base classes, CoCoOp could show lower accuracy on novel classes due to overfitting.